


Fayoum University
Faculty Of Engineering
Electronics & Communication Engineering



Project of

32-Bit Pipelined Mips Processor

Supervision by:

Prof. Dr. Gihan Naguib

Associate Professor, Department of Electrical Engineering

Eng. Jihad Awad

Teaching Assistant, Department of Electrical Engineering

Group Members:

- Omar Ahmed Othman
- Ahmed Osama Abdelghaffar
- Mohamed Nady Mahmoud

Department of Electronics and Communication Engineering (ECE)

Phase 1: Single Cycle Processor		
1	Introduction	5
2	Single Cycle Architecture	8
3	Register File	10
4	Arithmetic & Logic Unit (ALU)	12
5	Instruction Memory	19
6	Data Memory	21
7	Program Counter (PC)	24
8	Control Units	27
9	Main Control Unit	28
10	ALU Control Unit	32
11	Branch Control Unit	37
12	Full View of the Single-Cycle Processor	40
13	Simulation and Testing	41
Phase 2: Pipelined Processor		
1	Introduction	44
2	Pipeline Stages	45
3	Instruction Fetch (IF/ID)	46
4	Instruction Decode (ID/EXE)	48
5	Extender	51
6	Execute (EXE/MEM)	54
7	Memory Access (MEM/WB)	57
8	Write Back	61
9	Hazard Detect Forward and Stall Unit	62
10	PC Control Unit	67
11	Full view of Pipelined Processor	69
12	Team Work	70

Phase 1

Single Cycle Processor

Introduction

This project aims to design and implement a 32-bit RISC (Reduced Instruction Set Computer) processor based on the MIPS architecture using the Logisim digital circuit simulator. The processor supports a wide range of arithmetic, logic, memory, and control-flow instructions and is constructed following a structured and modular approach.

The primary objective of Phase 1 is to build a fully functional single-cycle processor, where each instruction is executed in a single clock cycle. This includes developing the complete Datapath, control unit, register file, ALU, instruction memory, and data memory, all integrated to support the full instruction set as defined in the project specification.

Instruction Set Architecture:

The design adheres to the custom ISA provided, which includes:

- 31 general-purpose registers (R1 to R31), with R0 hardwired to 0.
- Three instruction formats: R-type, I-type, and SB-type.
- Support for 32-bit word-addressable memory.
- Comprehensive instruction coverage for arithmetic, logic, shift, memory access, and branching.

Instructions Format:

R-Type Format

6-bit opcode (Op), 5-bit destination register number d, and two 5-bit source registers numbers S1 & S2 and 11-bit function field F

F₁₁ **S₂₅** **S₁₅** **D₅** **Op₆**

I-Type Format

6-bit opcode (Op), 5-bit destination register number d, and 5-bit source registers number S1 and 16-bit immediate (Imm16)

Imm₁₆ **S₁₅** **D₅** **Op₆**

SB-Type Format

6-bit opcode (Op), 5-bit register numbers (**S1**, and **S2**) and 16-bit immediate split into ({ImmU (11-bit) and ImmL(5-bit)})

ImmU₁₁ **S₂₅** **S₁₅** **ImmL₅** **Op₆**

This phase establishes the core functional blocks of the processor and serves as the foundation for the pipelined version in the next stage. The focus is on functional correctness, ensuring that all components interact properly and each instruction executes according to the defined semantics. Thorough testing is conducted through custom programs to validate correct execution and proper memory/register updates.

By completing Phase 1, we build a deep understanding of instruction execution, control signal generation, and Datapath behavior—providing a solid base for introducing pipeline stages, forwarding, and hazard handling in Phase 2.

Single-Cycle Architecture

The single-cycle architecture is a processor design in which each instruction is executed in one complete clock cycle. All stages of instruction execution including instruction fetch, decode, execution, memory access, and write-back, are performed in a single cycle. This design is straightforward and ideal for understanding the core operation of a processor, though it is not optimized for performance.

Datapath Overview:

The Datapath in the single-cycle MIPS processor consists of the following main components:

- **Register File:** Contains 32 registers (R0 to R31), with R0 hardwired to zero. Supports two read ports and one write port.
- **ALU:** A 32-bit Arithmetic and Logic Unit capable of performing arithmetic operations, bitwise logic, shifts, comparisons, and multiplication based on control signals.
- **Program Counter (PC):** A 20-bit register that holds the address of the current instruction.
- **Instruction Memory (IM):** Stores the machine code instructions; the PC value is used to fetch instructions.
- **Sign/Zero Extenders:** Extend the 16-bit immediate fields to 32-bit values for I-type and SB-type instructions.

- **Multiplexers (MUXes):** Used to select between different data inputs depending on the instruction type.
- **Data Memory (DM):** Used for load (LW) and store (SW) instructions to access 32-bit data words.
- **Control Unit:** Decodes the opcode and function fields to generate all necessary control signals for the datapath.


Instruction Execution

The datapath is designed to support execution of all instruction formats defined in the project:

- **R-Type:** The ALU performs operations using two source registers, and the result is written to a destination register.
- **I-Type:** Operations use one register and an immediate value; includes arithmetic, logical, and memory operations (LW).
- **SB-Type:** Used for branch and store instructions. The effective address is calculated using immediate values and source registers.

Register File

The register file is a core component of the single-cycle processor architecture, responsible for temporarily storing data during instruction execution. It consists of **32 general-purpose registers**, each 32 bits wide. These registers are used as operands for arithmetic/logical operations, for holding temporary values, and for storing results returned from the ALU or memory.



One of the registers, **R0**, is special—it is **hardwired to zero**. Reading from R0 always returns the value 0, and writing to it has no effect. This is useful in many operations, such as initializing values or representing "no result."


The register file supports the following features:

- **Two Read Ports:** To fetch the values of source operands (RS1 and RS2).
- **One Write Port:** To store the result into the destination, register (Rd).
- **Clocked Write:** The write operation occurs only on the rising edge of the clock signal and when the **Write Enable** signal is active.
- **Register Number Selection:** Each register is accessed using a 5-bit index (0–31).

Register File Structure

In our design, we have implemented a decoder equipped with a Write Register selector, facilitating the seamless designation of the target register for writing operations. This decoder selects from R1 to R31, serves as a pivotal component in directing data to the appropriate register within the Register File.


To ensure precise control over the writing operations, each terminal of the decoder is intricately linked with the RegWrite signal through AND Gates. This configuration allows for effective regulation of writing. Furthermore, we used two Multiplexers to enrich the functionality of our system by enabling the selection of output ports. This feature grants us to choose between utilizing one or both of the output ports depending on the instruction.




ALU (Arithmetic & Logic Unit)

The **Arithmetic Logic Unit (ALU)** is a fundamental component of the processor's datapath. It performs all arithmetic, logic, comparison, and shift operations required by the instruction set. The ALU receives two 32-bit input operands (from registers or immediate values), performs the operation specified by the control signals, and outputs the 32-bit result.

The ALU is designed to support **both R-type and I-type** instructions as defined in the instruction set architecture. Its functionality is controlled by the **ALU Control Unit**, which decodes the instruction's function field (F) or opcode and generates the appropriate control signals to select the desired operation.



1-Bit ALU Module



To construct the full 32-bit ALU, we designed a reusable **1-bit ALU building block**, which performs basic logical and arithmetic operations on a single pair of bits from the input operands. This modular design allows chaining 32 identical units to process 32-bit operands in parallel.

Inputs:

- A, B: 1-bit operands
- CarryIn: Incoming carry bit (for addition or subtraction)
- Operation [1:0]: Operation selector
- Ainvert: Determines if operand A should be inverted

- **Binvert:** Determines if operand B should be inverted (used for subtraction)

Outputs:

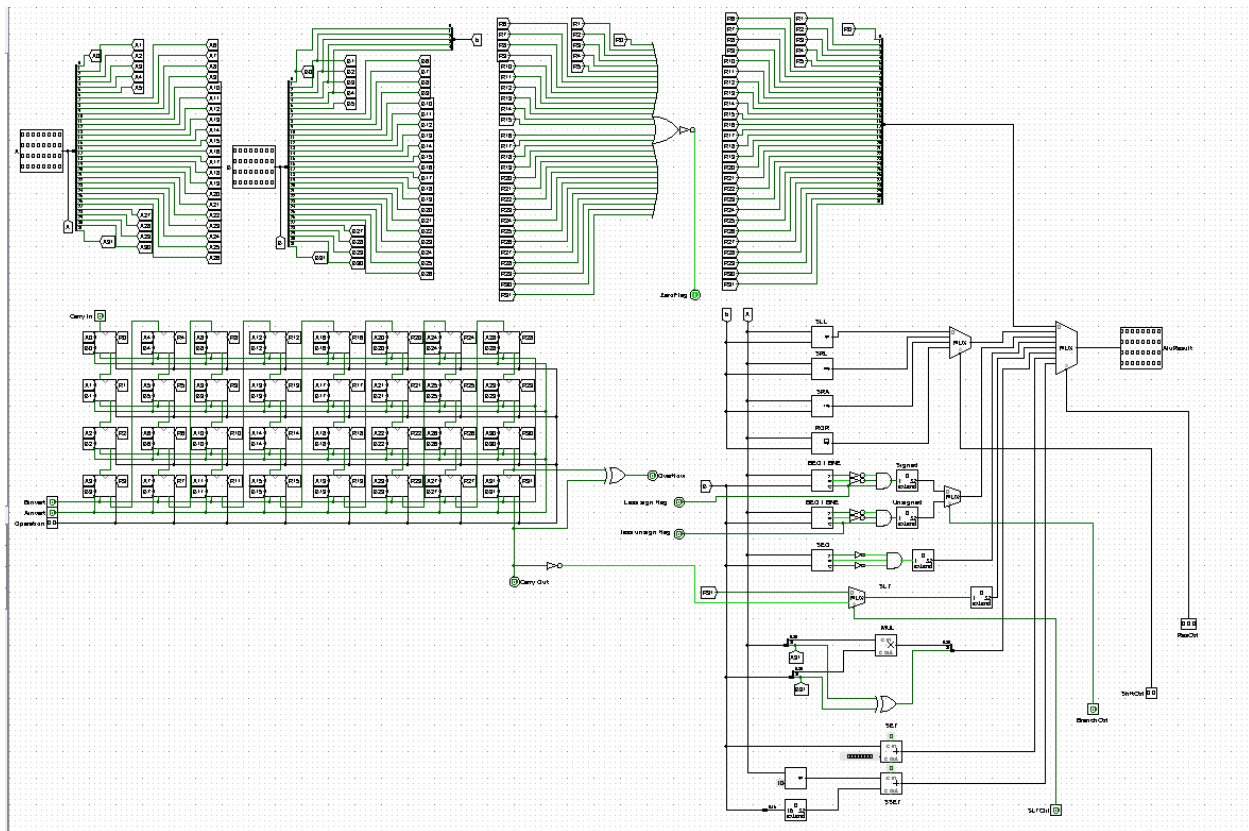
- **Result:** 1-bit result of the selected operation
- **CarryOut:** Carries out to the next higher bit

Supported Operations in Each 1-bit ALU:

- **AND:** Result = $A \& B$
- **OR:** Result = $A | B$
- **XOR:** Result = $A \wedge B$
- **NOR:** Result = $\sim(A \& B)$
- **ADD:** Result = $A + B + \text{CarryIn}$
- **SUB:** Performed as $A + (\sim B) + 1$ using Binvert and CarryIn

The full 32-bit ALU is constructed by **chaining 32 of these modules**.

32-Bit ALU Design



The 32-bit ALU is constructed by cascading 32 identical 1-bit ALU modules, allowing it to process two 32-bit operands in parallel. Each module is responsible for handling one bit of the operation, with carry and control signals propagated across the chain.

This modular structure simplifies design, improves maintainability, and makes it easier to implement arithmetic, logic, shift, and comparison operations.

Operation Handling

We have many types of operations like arithmetic, logical and shift& rotate operation, Let's discuss these types

1- Arithmetic Operations

- **Addition (ADD, ADDI, LW, SW, JALR)**
 - All 32 bits added using chained full-adders (1-bit ALUs).
 - Carry flows leftward through all bits.
- **Subtraction (SUB, SLT, SLTU, BEQ, BNE)**
 - Performed as $A + (\sim B) + 1$ using:
 - Binvert = 1
 - Cin = 1 at the LSB
 - Overflow and carry-out optionally detected at MSB.
- **Set-on-Less-Than (SLT, SLTU)**
 - Make Subtraction Operation and take the bit 31
 - At bit 31: Using it as a sign bit for SLT and use the Borrow flag for SLTU as Borrow = \sim CarryOut
- **Multiplication (MUL)**
 - A basic unsigned multiplier circuit or Logisim block used.
 - Outputs only the **lower 32 bits** of the 64-bit result.

2- Logical Operations

Handled bit-by-bit with no carry:

Instruction	Logic
AND, ANDI	$A \& B$
OR, ORI	$A B$
XOR, XORI	$A \wedge B$
NOR, NORI	$\sim(A B)$
SEQ, SEqi	$A == B ? 1 : 0$

- For immediate versions, **Operand B is zero-extended**.
- Comparison (SEQ) can be implemented by XOR \rightarrow NOT \rightarrow output 1 if equal

3 – Shift and Rotate Operations

These are handled using internal logic. Shift amount is typically taken from:

- **Lower 5 bits** of RS2 (R-type), or
- Immediate field (I-type).

Instruction	Operation
SLL, SLLI	$A \ll sa$
SRL, SRLI	$A \gg sa$ (logical)
SRA, SRAI	$A \gg sa$ (arithmetic; preserves sign bit)
ROR, RORI	A rotated right by sa bits

sa: shift amount

Our ALU is built in a modular and scalable way using chained 1-bit ALU units and dedicated logic for shift and rotate operations. By supporting a complete range of arithmetic, logical, and shift instructions

Its purely combinational design ensures immediate response to inputs, aligning perfectly with the single-cycle processor's requirement of completing each instruction in one clock cycle. As a final overview for some operations:

Supported Operations:

Operation	Description	Type
ADD	Adds two operands	Arithmetic
SUB	Subtracts B from A	Arithmetic
MUL	Multiplies A and B (lower 32 bits only)	Arithmetic
SLL	Logical shift left	Shift
SRL	Logical shift right	Shift
SRA	Arithmetic shift right	Shift
ROR	Rotate right	Shift
AND, OR, XOR, NOR	Bitwise logic operations	Logic
SLT, SLTU	Set if less than (signed/unsigned)	Comparison
SEQ	Set if equal	Comparison

Instruction Memory

The Instruction Memory (IM) is a read-only memory module responsible for storing the machine code of the program to be executed. Each instruction is 32 bits (4 bytes) wide and is word-aligned, meaning each address points to a full 32-bit word rather than individual bytes.

The Instruction Memory plays a key role in the fetch stage of the processor, where the instruction at the address held by the Program Counter (PC) is retrieved and passed to the rest of the datapath for decoding and execution.

Structure and Design

- **Addressing:**

The memory is word-addressable, meaning:

- Address 0 corresponds to instruction 0
- Address 1 corresponds to instruction 1
- and so on...

- **Size:**

The Instruction Memory is sized to store up to $2^{20} = 1,048,576$ instructions, consistent with a **20-bit PC**.

- **Access Type:**

The memory is **read-only** during program execution. Instructions are preloaded during initialization or simulation.

- **Input:**

- PC (20-bit): The address of the current instruction.

- **Output:**

- Instruction (32-bit): The 32-bit instruction at the address PC.


- **Instruction Alignment:**

Since instructions are word-aligned, the PC is incremented by **1** after each instruction fetch, not by 4

Instruction Fetching Logic

The Instruction Memory is directly accessed using the PC:

$$\text{Instruction} \leftarrow \text{IM}[\text{PC}]$$



After fetching, the instruction is split into fields:

- Opcode (6 bits)
- RS1, RS2=Rt, Rd = ImmediateL (5 bits each)
- Function (11 bits for R-type) = ImmediateU or Immediate (16 bits for I-type/SB-type)


Data Memory

The Data Memory (DM) is a critical component of the processor responsible for storing and retrieving data during program execution. It is a separate memory block from the Instruction Memory.

Data Memory is primarily used by load (LW) and store (SW) instructions, which operate on word-sized (32-bit) data.

Following the data memory, a multiplexer (MUX) is used to select between the ALU result and the memory output, depending on whether the instruction is a computational instruction (like ADD, SUB, SLT) or a memory access instruction (like LW). Additionally, another MUX is placed before writing to the register file, which selects between the normal data (from ALU or memory) and the link address ($PC + 1$) in the case of JALR instructions.

These MUXes play a critical role in routing the correct result back to the register file based on the instruction type and control signals.



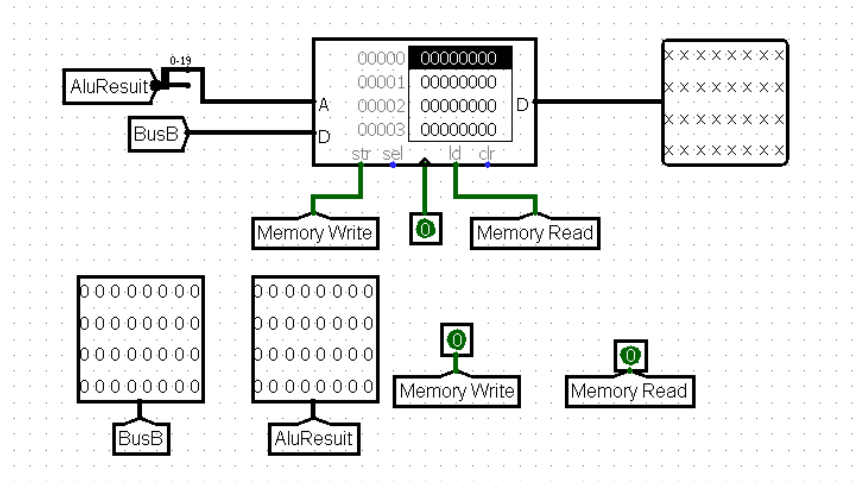
Structure and Design

- **Size:**

The memory is configured to hold $2^{20} = 1,048,576$ words (4 MB total if word-addressed). This matches the instruction memory's address range.

- **Word Addressable:**

Each memory location holds **one 32-bit word**, and the **addresses increment by 1** for each word. Byte addressing is not supported in this design.



- **Access Ports:**

- **Address** (32-bit): The effective memory address, typically the output of the ALU.
- **Data In** (32-bit): Data to be written into memory (SW) = Bus B.
- **MemWrite** (1-bit): Control signal to enable write operation.
- **MemRead** (1-bit): Control signal to enable read operation.
- **ReadData** (32-bit): Output of the memory (LW).

Operation Modes

1. Load Word (LW):

- Reads 32-bit word from Address = RS1 + sign extend (Imm16)
- Output is placed on the ReadData bus
- Used when MemRead = 1 and MemWrite = 0

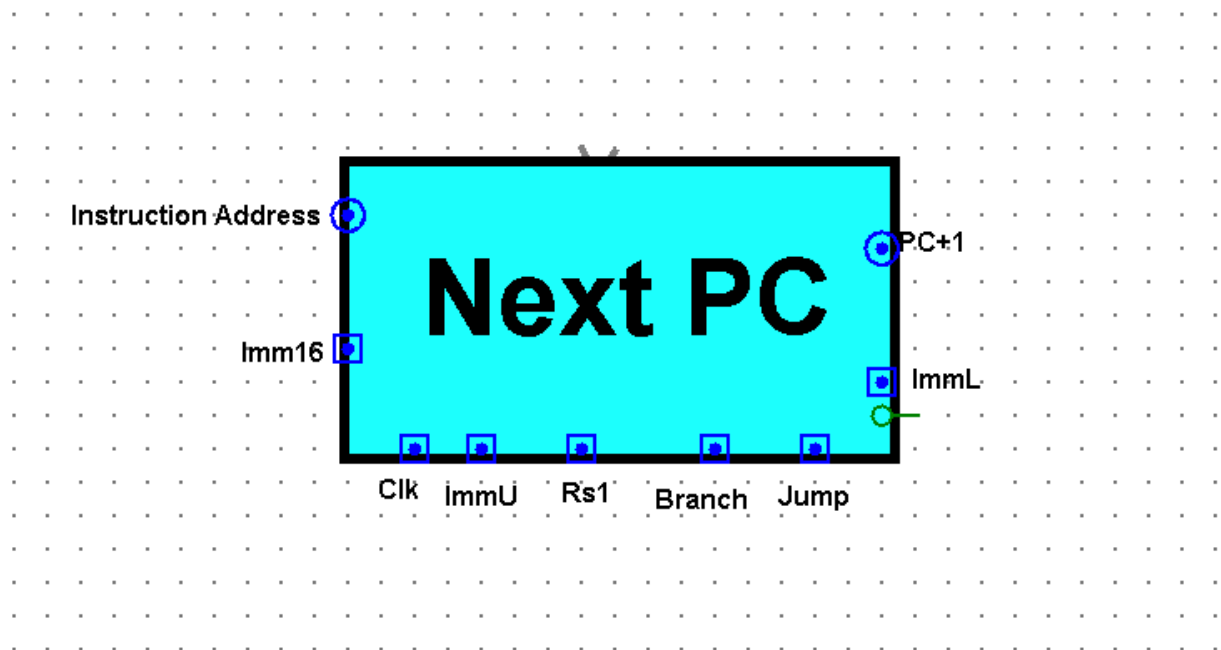
2. Store Word (SW):

- Writes 32-bit WriteData (RS2) to address RS1 + sign extend ({ImmU, ImmL})
- Occurs only when MemWrite = 1 and MemRead = 0

Program Counter (PC)

The Program Counter (PC) is a special-purpose 20-bit register that holds the address of the current instruction to be executed. It is an essential component in the instruction fetch stage of the processor and plays a central role in controlling the flow of program execution.

Since the instruction memory is word-addressable and each instruction is 32 bits (4 bytes), the PC increments by 1 after each instruction to point to the next word-aligned instruction in memory.



The Program Counter ensures the sequential flow of instruction execution, and its ability to redirect control to new instruction addresses allows for the implementation of loops, branches, and procedure calls. Together with the control unit and immediate extension logic, the PC enables both simple and complex control flows in the processor.

Structure and Operation

Inputs:

- PC: Current program counter value (20 bits).
- Imm16: 16-bit immediate value used for JALR.
- ImmU, ImmL: 11-bit and 5-bit parts of the 16-bit immediate used in SB-type branch instructions.
- RS1: Value from the source register for computing JALR jump addresses.
- Branch: Control signal that enables conditional branching.
- Jump: Control signal that enables jumps (like JALR).

Outputs:

- Instruction Address: Sent to Instruction Memory to fetch the next instruction.
- PC + 1: Default next instruction address (used for sequential execution).

the PC update logic considers three possible target addresses:

1. **Sequential Address:** PC + 1 — the next instruction in normal execution.
2. **Branch Address:** PC + sign extend ({ImmU, ImmL}) — used for conditional branches (BEQ, BNE, etc.).
3. **Jump Address:** RS1 + sign extend (Imm16) — used in indirect jumps (JALR and its pseudo-instructions).

Control Units

In a single-cycle MIPS processor, the control logic is responsible for generating all the necessary signals that guide the flow of data through the datapath. These signals determine:

- Which registers to read/write,
- Whether to perform an ALU operation or memory access,
- How to update the Program Counter (PC),
- And which operation the ALU should perform.

The control logic is typically split into two major components:

1. **Main Control Unit**

Generates high-level control signals based on the instruction opcode.

2. **ALU Control Unit**

Generates the specific operation code for the ALU based on the instruction function field and ALUOp signal from the Main CU.

3. **Branch Control Unit**

Determines whether we going to take the branch address or not according to the instruction

Main Control Unit

The Main Control Unit is the central block responsible for decoding the 6-bit instruction opcode and generating the high-level control signals that guide data flow throughout the processor. It does not determine the exact ALU operation (which is handled by the ALU Control Unit), but instead decides *what* the processor should do on a broader level — such as whether to write to a register, access memory, perform a jump, or select between immediate and register operands.



Inputs:

- Opcode (6 bits): The main identifier field of the instruction.

Outputs:

Signal	Purpose
RegWr	Enables writing to the register file. Active for most R-type and I-type instructions, and JALR.
MemRd	Enables reading from Data Memory (LW).
MemWr	Enables writing to Data Memory (SW).
MemtoReg	Selects data source to write back to register: from memory (LW) or ALU result .
SW	Active only for store word (SW) instruction — used to differentiate control logic specific to stores.
SSET	Active for SSET instruction , used to select special write-back logic: $Rd = \{Rd [15:0], Imm16\}$.
ALUsrc	Selects second ALU operand: 0 for register, 1 for immediate.
ExtOp	Controls how the immediate is extended:

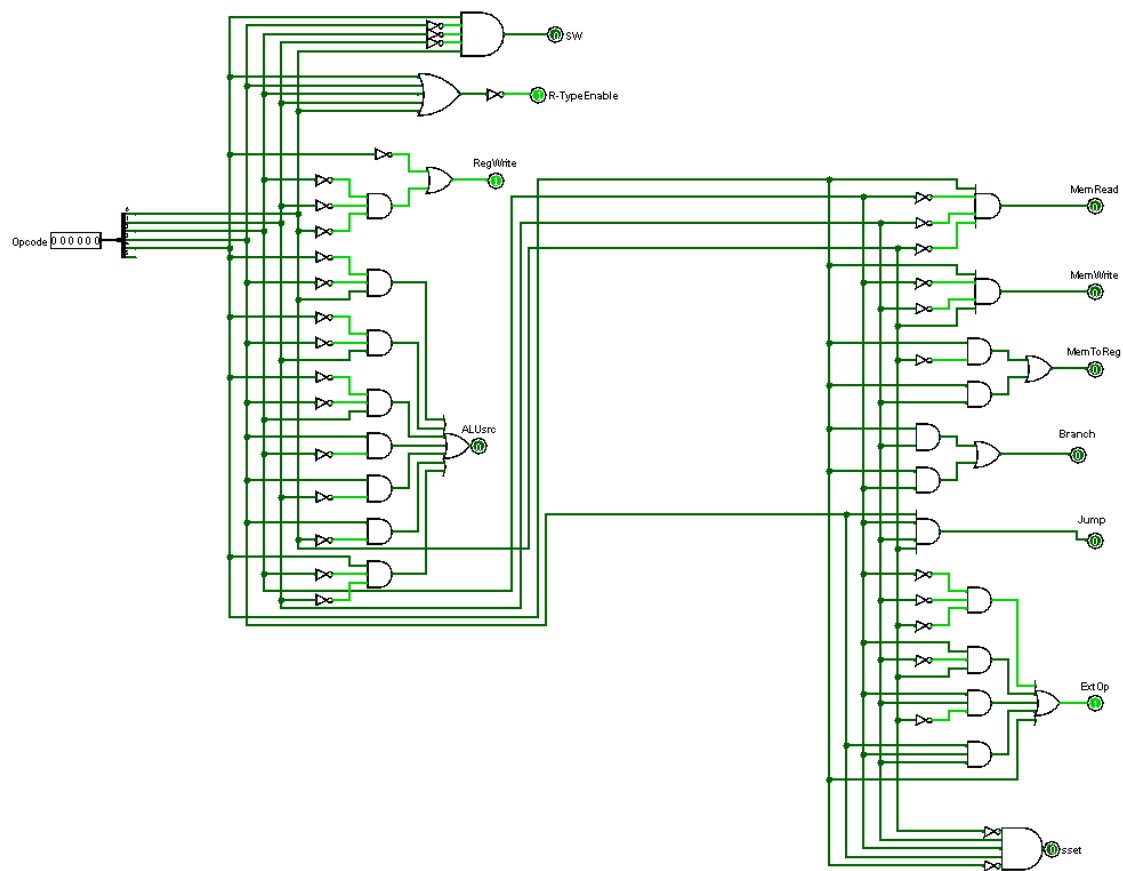
- 0: Zero extension (e.g., ANDI, ORI),
- 1: Sign extension (e.g., ADDI, SLTI).

- **R-type Enable:**

Set high when the instruction is R-type (Opcode = 0). Used to enable function field decoding in ALU CU.

- **Jump:** High for JALR instruction or its pseudo instructions (JR, J, etc.).
- **Branch:** High for SB-type branch instructions (BEQ, BNE, etc.) to activate Branch Control Unit.

It interprets the **opcode** field of each instruction and generates the appropriate **control signals** that drive the datapath components. As shown in the circuit above, this unit consists primarily of logic gates that decode the 6-bit opcode to produce the necessary output signals.



Main Control Unit Truth Table:

Opcode	RegWr	RegWr	MemRd	MemWr	MemtoReg	Branch	JumpAL	ExtOp	R-TypeEn	SSET	SW
0	1	0	0	0	0	0	0	x	1	x	x
1	1	1	0	0	0	0	0	0	0	x	0
2	1	1	0	0	0	0	0	0	0	x	0
3	1	1	0	0	0	0	0	0	0	x	0
4	1	1	0	0	0	0	0	0	0	x	0
5	1	1	0	0	0	0	0	1	0	x	0
6	1	1	0	0	0	0	0	1	0	x	0
7	1	1	0	0	0	0	0	0	0	x	0
8	1	1	0	0	0	0	0	1	0	x	0
9	1	1	0	0	0	0	0	0	0	x	0
10	1	1	0	0	0	0	0	0	0	x	0
11	1	1	0	0	0	0	0	0	0	x	0
12	1	1	0	0	0	0	0	0	0	x	0
13	1	1	0	0	0	0	0	1	0	x	0
14	1	1	0	0	0	0	0	1	0	1	0
15	1	0	0	0	0	0	1	1	0	x	0
16	1	1	1	0	1	0	0	1	0	x	0
17	x	1	0	1	0	0	0	1	0	x	1
18	x	0	0	0	x	1	0	1	0	x	0
19	x	0	0	0	x	1	0	1	0	x	0
20	x	0	0	0	x	1	0	1	0	x	0
21	x	0	0	0	x	1	0	1	0	x	0
22	x	0	0	0	x	1	0	1	0	x	0
23	x	0	0	0	x	1	0	1	0	x	0

Special Case: Opcode = 000000 (R-type, SSET, and SW)

When the Opcode is 000000, the instruction is classified as a special type — commonly referred to as R-type, but not limited to arithmetic and logic operations. In this case, the function field (func) of the instruction is used to further decode the exact operation in the ALU Control Unit or handle unique behaviors in the Main Control Unit.

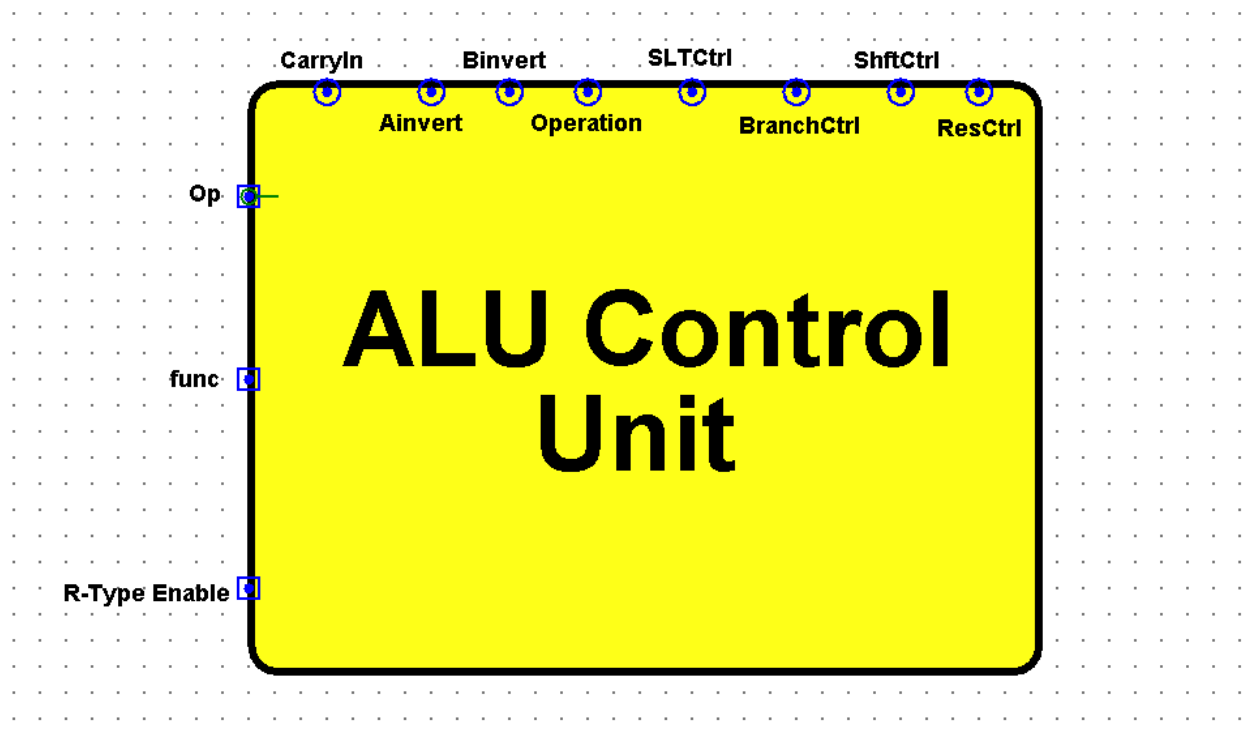
Under Opcode = 0, the processor supports:

- R-type operations like ADD, SUB, AND, OR, SLT, SLL, SRL, etc.
- SSET instruction (Special immediate move to upper half of register), enabled by asserting the SSET signal.
- SW (Store Word), which also uses Opcode = 0 in this implementation, distinguished via its func field and activates the SW and MemWr control lines

ALU Control Unit

The ALU Control Unit is a key component in the execution stage of the processor's datapath. Its primary function is to generate the appropriate control signals that configure the Arithmetic Logic Unit (ALU) to perform the desired operation—whether arithmetic, logical, or shift-related.

While the Main Control Unit classifies the instruction type based on the opcode (e.g., R-type, I-type, branch), it is the responsibility of the ALU Control Unit to decode the instruction further and select the exact ALU operation. This distinction ensures a modular and scalable control logic architecture.



Functionality and Input Signals

The ALU CU takes the following key inputs:

- **Opcod**: A 6-bit field used to distinguish high-level instruction types.
- **Func**: A 6-bit function code used specifically by R-type instructions to determine the exact ALU operation.
- **R-typeEnable**: A control flag from the Main CU that signals whether the instruction is R-type.

Based on these inputs:

- If R-typeEnable is 1, the ALU CU interprets the Func field to determine the ALU operation (e.g., ADD, SUB, SLT, SLL, etc.).
- If R-typeEnable is 0, the ALU operation is derived directly from the Opcode, as in immediate-type instructions like ADDI, ANDI, or ORI.

Generated Control Signals

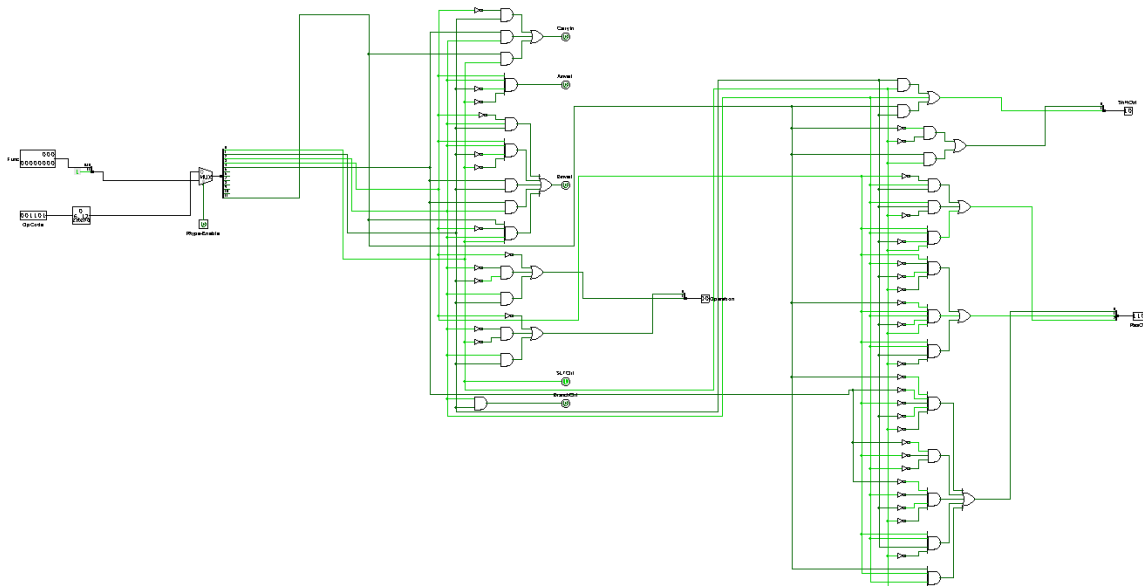
Signal	Purpose
CarryIn	Controls the initial carry bit input; crucial for operations like subtraction.
Ainvert	Inverts the A operand when needed (e.g., for NAND or subtraction).
Binvert	Inverts the B operand; essential for subtract and logical NOR.
Operation	A multi-bit signal that selects the arithmetic or logic operation (e.g., ADD, AND, OR).
SLTCtrl	Select between whether the operation is signed or unsigned.
BranchCtrl	Indicates if the operation affects control flow (e.g., BEQ, BNE, BGT).
ShftCtrl	Activates shift operations (e.g., logical shift left/right).
ResCtrl	Selects the final result source (e.g., ALU output, shift unit output, SLT).

The unit is implemented using pure combinational logic:

- If **R-Type Enable** is active, the control logic depends entirely on the func field to determine the operation.
- Otherwise, the logic uses the opcode to determine the ALU behavior for I-type instructions.
- Logical gates (AND, OR, NOT) are used to match specific bit patterns in the opcode/func to assert the correct control signals.
- This setup ensures only one active control combination for a given instruction.

Example Logic:

- If R-TypeEnable = 1 and func = 100000 → ADD → Operation = 10, Ainvert = 0, Binvert = 0.
- If opcode = 001000 (ADDI) → same ALU behavior but triggered through opcode logic instead of func.



ALU Control Unit Truth Table

At R-Type Enable = 0

Opcode	CarryIn	Ainvert	Binvert	Operation	SLTCtrl	BCtrl	ShftCtrl	ResCtrl
1	x	0	0	xx	x	x	00	001
2	x	0	0	xx	x	x	01	001
3	x	0	0	xx	x	x	10	001
4	x	0	0	xx	x	x	11	001
5	0	0	0	11	x	x	xx	000
6	1	0	1	11	0	x	xx	100
7	x	0	1	11	1	x	xx	100
8	x	0	0	xx	x	x	xx	011
9	x	0	0	10	x	x	xx	000
10	x	0	0	1	x	x	xx	000
11	x	0	0	0	x	x	xx	000
12	x	1	1	0	x	x	xx	000
13	x	0	0	xx	x	x	xx	110
14	x	0	0	xx	x	x	xx	111
15	0	0	0	11	x	x	xx	000
16	0	0	0	11	x	x	xx	000
17	0	0	0	11	x	x	xx	000
18	1	0	1	11	x	0	xx	xxx
19	1	0	1	11	x	0	xx	xxx
20	1	0	1	11	x	0	xx	xxx
21	1	0	1	11	x	0	xx	xxx
22	1	0	1	11	x	1	xx	xxx
23	1	0	1	11	x	1	xx	xxx

ALU operation is derived directly from the Opcode

At R-Type Enable = 1

Func	CarryIn	Ainvert	Binvert	Operation	SLTCtrl	BCtrl	ShftCtrl	ResCtrl
0	x	0	0	xx	x	x	00	001
1	x	0	0	xx	x	x	01	001
2	x	0	0	xx	x	x	10	001
3	x	0	0	xx	x	x	11	001
4	0	0	0	11	x	x	xx	000
5	1	0	1	11	x	x	xx	000
6	1	0	1	11	0	x	xx	100
7	1	0	1	11	1	x	xx	100
8	x	0	0	xx	x	x	xx	011
9	x	0	0	10	x	x	xx	000
10	x	0	0	01	x	x	xx	000
11	x	0	0	00	x	x	xx	000
12	x	1	1	00	x	x	xx	000
13	x	0	0	xx	x	x	xx	101

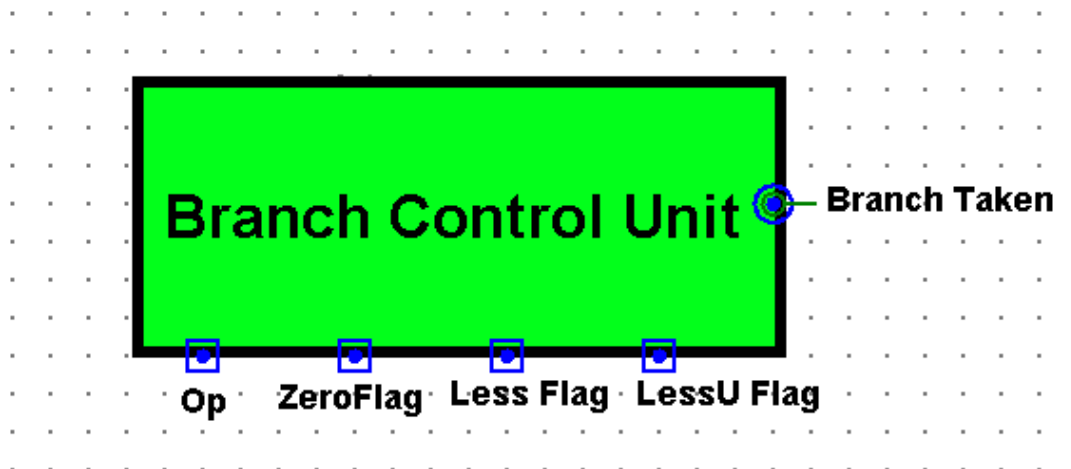
ALU CU interprets the Func field to determine the ALU operation (e.g., ADD, SUB, SLT, SLL, etc.).

Branch Control Unit

The Branch Control Unit (BCU) is a dedicated logic module that evaluates branch decisions based on a limited yet sufficient set of flags provided by the ALU. This unit directly determines whether control should be transferred to a branch target address, or if execution should continue sequentially.

Inputs:

- Opcode: 6 bits from instruction
- ALU Flags:
 - Zero: Set when the two operands are equal.
 - Less (signed): Set when $rs1 < rs2$ in signed comparison.
 - LessU (unsigned): Set when $rs1 < rs2$ in unsigned comparison.



Branch Instructions:

Instruction	Meaning	Condition Evaluated	Logic Expression
BEQ	Branch if Equal	Zero == 1	Take branch if equal
BNE	Branch if Not Equal	Zero == 0	Take branch if not equal
BLT	Branch if Less Than	Less == 1 (signed comparison)	Take branch if $rs1 < rs2$
BGE	Branch if Greater/Equal	Less == 0 (signed comparison)	Take branch if $rs1 \geq rs2$
BLTU	Branch if Less Than (U)	LessU == 1 (unsigned comparison)	Take branch if $rs1 < rs2$ (U)
BGEU	Branch if GE (Unsigned)	LessU == 0	Take branch if $rs1 \geq rs2$ (U)

All branch decisions are binary (take/not take) and are based only on one of the three flags per instruction.


Output:

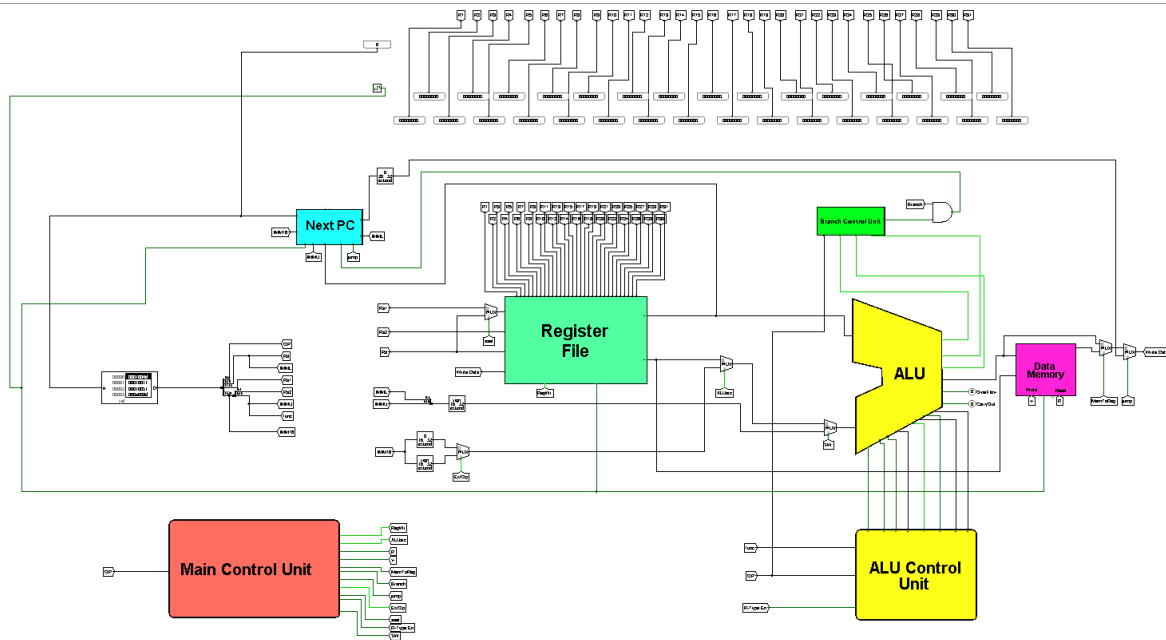
- BranchTaken (1-bit signal): Activated when the condition for the current branch instruction is met. Then Anded with Branch Signal from Main CU.

This signal resulted from the AND Operation feeds into the Program Counter (PC) control logic, enabling it to select between $PC + 1$ and the branch target address.

Branch Control Unit Truth Table

Opcode	Instruction	Zero	Less	LessU	BranchTaken
010010 (18)	BEQ	1	X	X	1
010011 (19)	BNE	0	X	X	1
010100 (20)	BLT	X	1	X	1
010101 (21)	BGE	X	0	X	1
010110 (22)	BLTU	X	X	1	1
010111 (23)	BGEU	X	X	0	1





After discussing all individual components in detail — including the register file, ALU, instruction and data memory, control units, and branching logic, we now present the **full single-cycle datapath**. This unified structure represents how the processor executes an instruction from fetch to write-back within a single clock cycle. At a glance, we can observe how control signals orchestrate data flow between components, how multiplexers guide conditional paths (such as jumps or branches), and how immediate values are extended and routed appropriately. The datapath captures the harmony between computation, memory access, and control decisions. This integrated view helps visualize the complete instruction lifecycle and emphasizes the processor's one-cycle-per-instruction design philosophy.

Simulation and Testing

The table below summarizes the executed instructions, their hexadecimal machine code, and the expected output. These outputs were compared with the actual values from simulation waveforms, and results showed full match with the expected behavior.

Instruction	Hex Representation	Expected Output
SET R1, 0x0001	0001004D	R1 = 0x00000001
SW R1, 0(R0)	00010011	Mem [0] = 0x00000001
SW R1, 1(R0)	00010051	Mem [1] = 0x00000001
SET R2, 0x000A	000A008D	R2 = 0x0000000A
SW R2, 2(R0)	00020091	Mem [2] = 0x0000000A
SET R3, 0x1289	128900CD	R3 = 0x00001289
SSET R3, 0x45AC	45AC00CE	R3 = 0x128945AC
SW R3, 60(R0)	00230711	Mem [60] = 0x128945AC
SET R4, 0x0500	0500010D	R4 = 0x00000500
SSET R4, 0x7342	7342010E	R4 = 0x05007342
SW R4, 61(R0)	00240751	Mem [61] = 0x05007342
SET R1, 0x0384	0384004D	R1 = 0x00000384
SET R8, 0x1234	1234020D	R8 = 0x00001234
SSET R8, 0x5678	5678020E	R8 = 0x12345678
ADDI R5, R1, 20	00140945	R5 = 0x398
XOR R3, R1, R5	012508C0	R3 = 0x1C
ADD R4, R8, R3	00834100	R4 = 0x12345694
LW R1, 0(R0)	00000050	R1 = 0x00000001
LW R2, 1(R0)	00010090	R2 = 0x00000001
LW R3, 2(R0)	000200D0	R3 = 0x0000000A

SUB R4, R4, R4	00A42100	R4 = 0x0
ADD R4, R2, R4	00841100	R4 += R2
SLT R6, R2, R3	00C31180	R6 = (R2 < R3)? 1: 0
BEQ R6, R0, done	000030D2	Branch if R6 == 0
ADD R2, R1, R2	00820880	R2 += R1
BEQ R0, R0, Loop1	FFE00712	Unconditional branch (loop)
SW R4, 0(R0)	00040011	Mem [0] = R4 = 0x37
MUL R10, R2, R3	01A31280	R10 = R2 * R3 = 0x64
SRL R14, R10, R4	00245380	R14 = R10 >> R4 (logical) = 0
SRA R15, R10, R4	004453C0	R15 = R10 >> R4 (arith) = 0
RORI R26, R14, 5	00057684	R26 = ROR (R14, 5) = 0
JALR R7, R0, func	002501CF	Jump to func, store return address in R7
SET R9, 0x4545	4545024D	R9 = 0x00004545
SET R10, 0x4545	4545028D	R10 = 0x00004545
BGE R10, R9, L1	00095095	Branch if R10 >= R9 (taken)
ANDI R23, R1, 0xFFFF	FFFF0DCB	Skipped
BEQ R0, R0, L1	00000012	Infinite loop (halt)
OR R5, R2, R3	01431140	R5 = R2
LW R1, 0(R0)	00000050	R1 = 0x37
LW R2, 5(R1)	00050890	R2 = Mem [R1 + 5] = 0x128945AC
LW R3, 6(R1)	000608D0	R3 = Mem [R1 + 6] = 0x05007342
AND R4, R2, R3	01631100	R4 = R2 & R3 = 0x4100
SW R4, 0(R0)	00040011	Mem [0] = R4
JALR R0, R7, 0	0000380F	Return to caller (JR R7)

Phase 2

Pipelined Processor

Introduction

In this phase of the project, we took our working single-cycle MIPS processor and transformed it into a pipelined version to improve performance. Instead of executing one instruction at a time from start to finish, pipelining allows multiple instructions to be processed at the same time by breaking the execution into smaller stages. This means while one instruction is being decoded, another can be fetched, and yet another can be executed—just like an assembly line.

To make this work, we divided the processor into five main stages: Instruction Fetch, Decode, Execute, Memory Access, and Write Back. We added pipeline registers between these stages to keep track of each instruction as it moves through the pipeline.

Of course, pipelining brings its own challenges. We had to handle data hazards, which happen when instructions depend on each other. To deal with this, we implemented forwarding logic so that the needed values are passed along without waiting. We also added stalling logic to pause the pipeline when a load instruction is immediately followed by one that needs its result. Control hazards, caused by branches and jumps, were also addressed to keep the processor on track.

Overall, this phase was about making the processor faster and smarter, while still correctly handling all the instructions. It was a big step forward in turning our design into something that works more like a real-world CPU.

Pipeline Stages

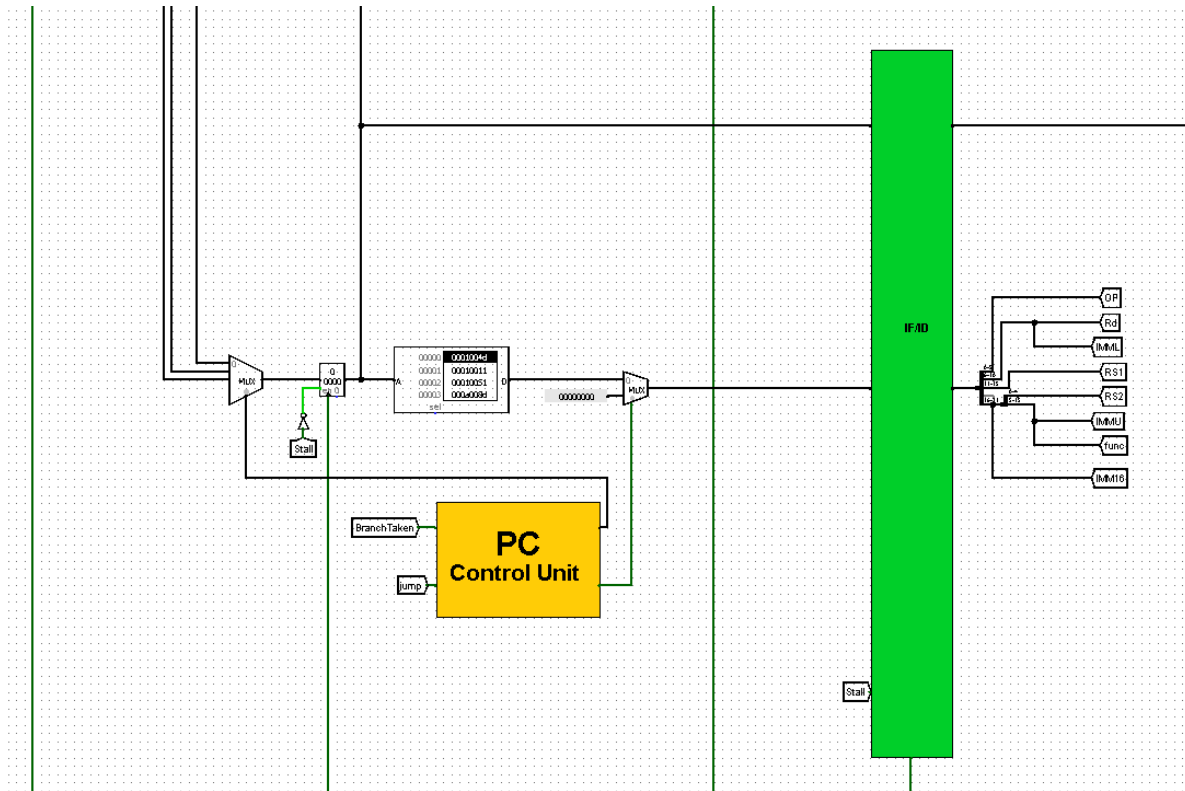
Our pipelined processor is based on the classic 5-stage RISC pipeline architecture, which is a well-established structure used in many modern CPUs. The idea behind pipelining is similar to an assembly line: instead of waiting for one instruction to finish completely before starting the next, we split the work into stages, so multiple instructions can be in progress at the same time—each at a different step.

Every instruction goes through the same five stages:

- Instruction Fetch (IF)
- Instruction Decode (ID)
- Execute (EX)
- Memory Access (MEM)
- Write Back (WB)

We will discuss every stage later.

Instruction Fetch



The Instruction Fetch stage is the first step in the processor's execution pipeline. Its main job is to retrieve the next instruction to be executed from the Instruction Memory, using the address stored in the Program Counter (PC). This stage sets the rhythm for the rest of the pipeline, as it continuously feeds the processor with instructions.

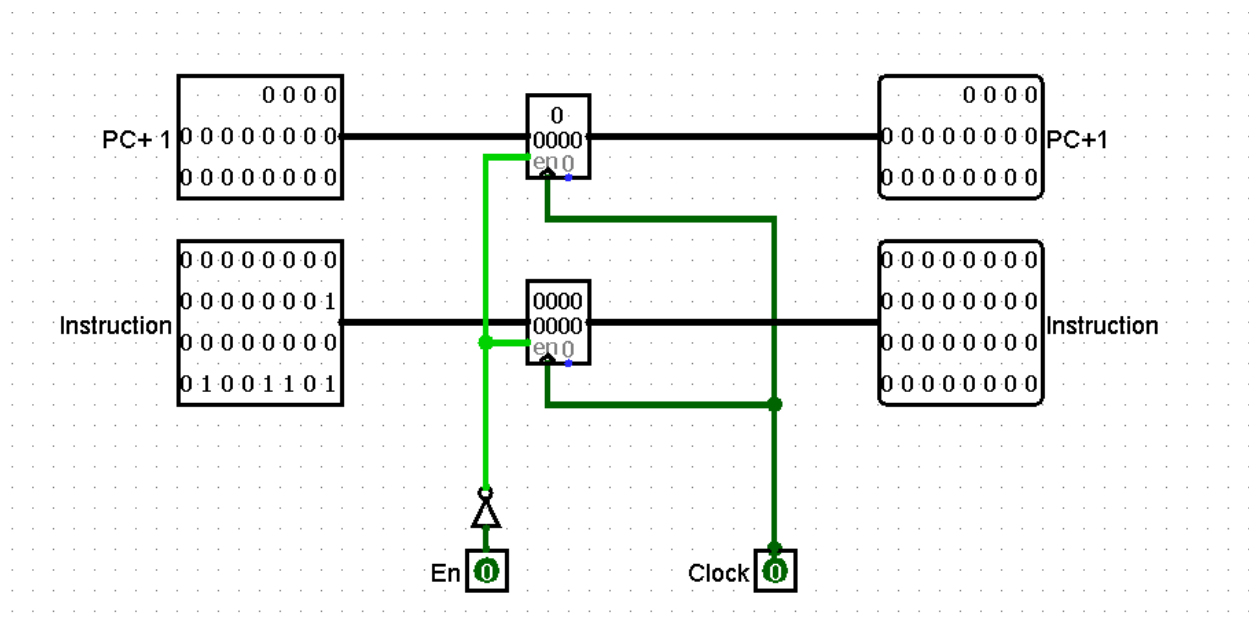
In our pipelined processor, the PC holds the address of the current instruction and is updated every cycle, unless a branch or jump alters the control flow. Since our instruction memory is word-addressable, and each instruction is 32 bits (4 bytes), we increment the PC by 1 to move to the next instruction.

Once the instruction is fetched, it is stored in the IF/ID pipeline register, along with the current PC value. This allows the next stage (Instruction Decode) to begin processing the instruction while the IF stage continues fetching the following one.

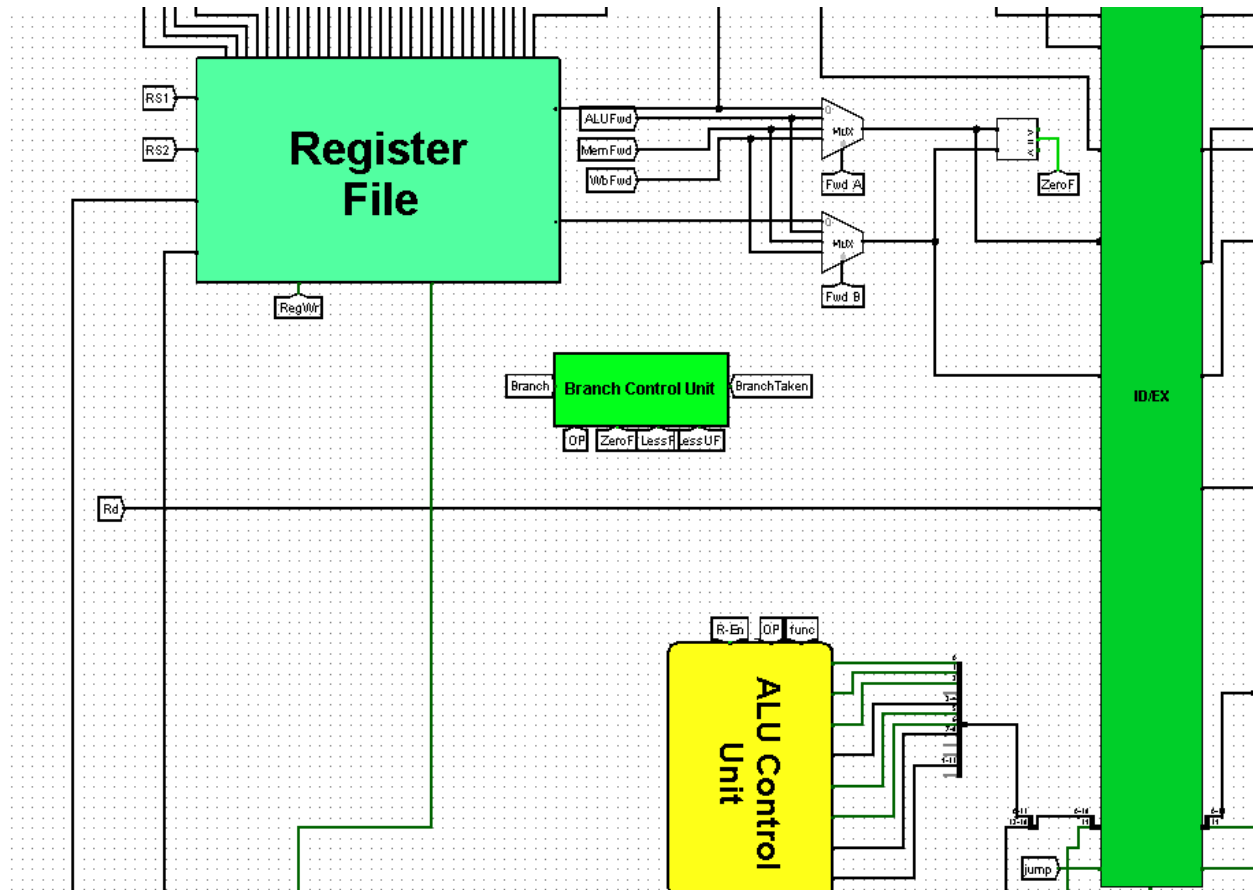
The simplicity of this stage hides its importance—it keeps the pipeline flowing smoothly. In branch or jump situations, the PC may be updated with a new address based on the outcome of the previous instruction. In such cases, the IF stage is also responsible for flushing or replacing incorrectly fetched instructions to ensure correct program execution.

This Enable Signal will be later the Stall Signal.

We will discuss later in the Hazard unit.



Instruction Decode



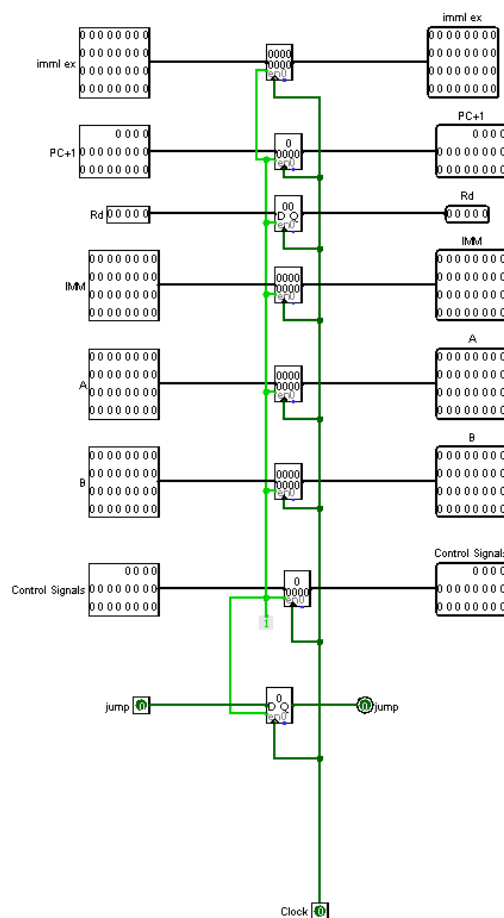
The Instruction Decode stage is the second step in the pipeline and plays a crucial role in understanding what the fetched instruction actually means.

Once an instruction is fetched from memory, it enters the ID stage where it is decoded into its components—operation type, source and destination registers, and immediate values if any.

During this stage, the control unit interprets the opcode and generates the necessary control signals that will guide the instruction through the rest of the pipeline. For R-type instructions, the decode logic identifies the two source

registers and the destination register. For I-type and SB-type instructions, it also extracts and processes the immediate value—either by sign-extending or zero-extending it to 32 bits based on the instruction.

A key task in this stage is reading the values from the register file. The values of the source registers (RS1 and RS2) are read and forwarded to the next stage along with the instruction and any calculated control signals. Meanwhile, the destination register is noted for later use in the Write Back stage.



Stored Signals

Label	Meaning
imm1 ex	sign-extended immediate (used for I-type instructions)
PC+1	Next program counter (could be PC+4); used in branch calculations
Rd	Destination registers field (usually for R-type instructions)
IMM	The raw immediate value and further extended value
A	Value read from register rs (operand 1)
B	Value read from register rt (operand 2)
Control Signals	Encoded control lines (e.g., RegDst, ALUOp, MemRead, MemWrite, etc.)
jump	flag indicating a jump instruction is in progress (JALR)


Once the EX-stage begins:

- A and B go to the ALU for computation.
- imm1 ex or IMM can be used as an ALU operand via a mux (depending on whether the instruction is R-type or I-type).
- Rd helps determine the destination register (used with RegDst).
- PC+1 can be used for branch target calculation (usually with imm1 ex).
- Control Signals determine what happens downstream (e.g., if the result goes to memory, register, etc.).
- jump helps decide PC redirection in case of JALR instruction.

We also added a new block called it **Extender**

Extender

This block handles all types of **immediate value extensions** and **address calculations** needed for the next pipeline stage (ID/EX).



In our pipelined processor, a specialized Extender block was implemented as part of the Instruction Decode (ID) stage. Its purpose is to generate properly formatted and extended values for branch targets, jump addresses, and immediate operands required by the ALU in later stages.

This unit takes the raw immediate fields from the instruction and processes them as follows:

Branch Address Calculation

For SB-type branch instructions (e.g., BEQ, BNE, etc.), the 16-bit immediate is split into ImmU (upper 11 bits) and ImmL (lower 5 bits). The Extender block:

- Concatenates them into a 16-bit value.
- Sign-extends this to 20 bits.
- Adds it to PC + 1 to compute the target branch address.

Jump Address Calculation

For JALR (Jump and Link Register) instruction:

- The Extender block takes RS1 and a sign-extended 16-bit immediate.
- It adds them together to form the effective jump address.

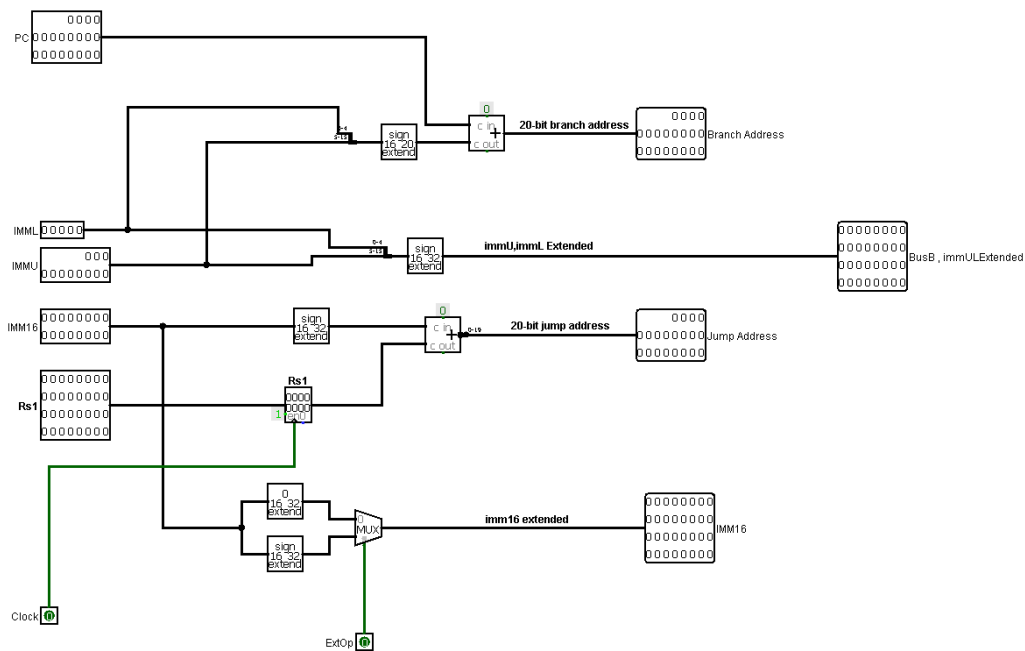
Immediate Extension for ALU Instructions

Depending on the type of I-format instruction:

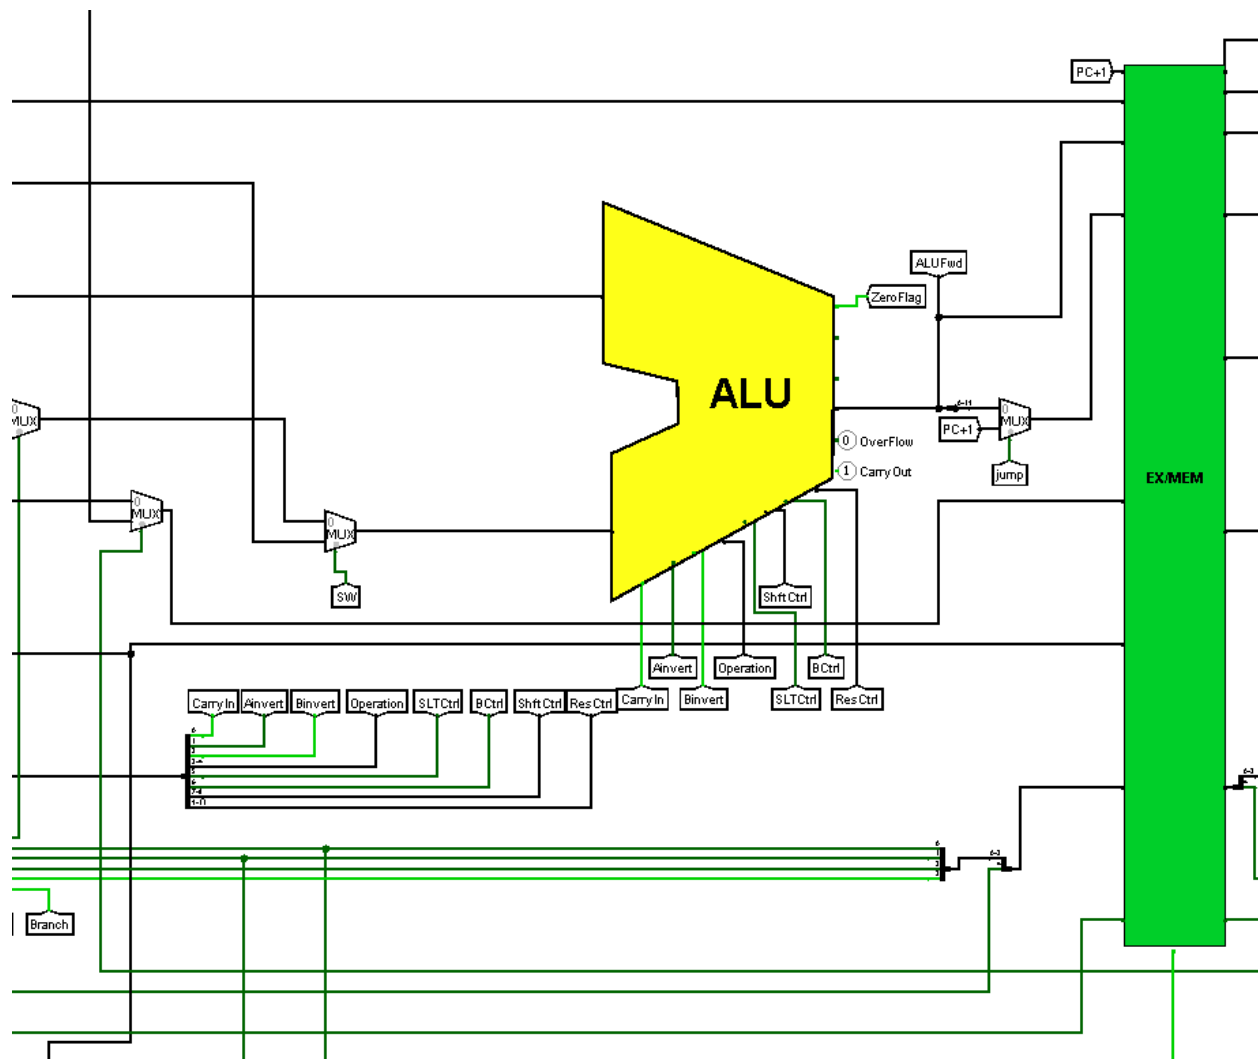
- The 16-bit immediate may be either sign-extended (for arithmetic and branch operations) or zero-extended (for bitwise operations like ANDI, ORI, etc.).
- A control signal (ExtOp) selects between sign-extension and zero-extension via a MUX.

- The resulting 32-bit extended value is passed to the ID/EX pipeline register as imm16_extended.

This modular design helps reduce complexity in the datapath and ensures that all extended values are available and correctly formatted before moving to the Execute (EX) stage. By isolating the logic in one place, it also simplifies debugging and improves modularity of the decode logic.



Execute



The Execute stage is where the processor performs actual computations—whether arithmetic, logic, address calculation, or comparisons—based on the decoded instruction and its operands. Once these operations are complete, the results, along with other necessary signals, must be preserved and passed to the MEM stage through the EX/MEM pipeline register.

This pipeline register acts as a snapshot of all the outputs from the EX-stage. Without it, the next stage wouldn't know what result to use, what memory

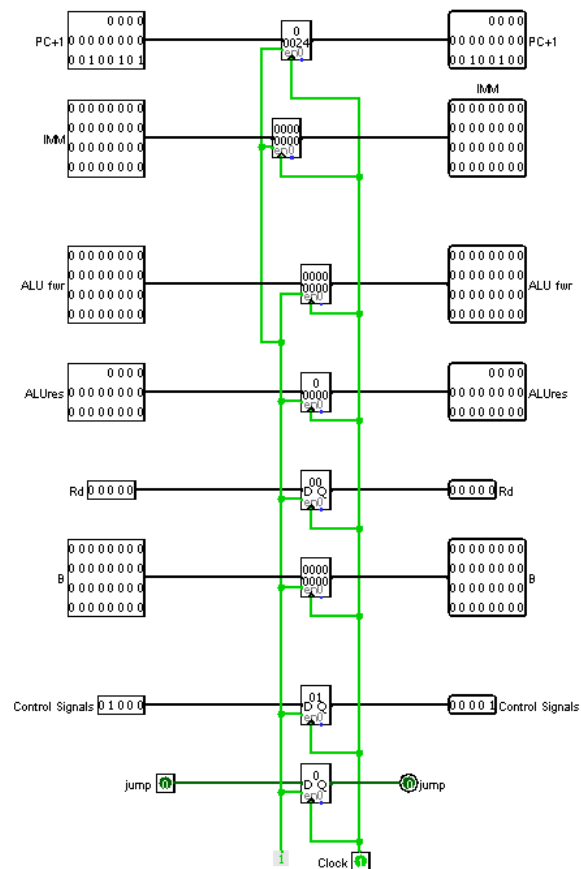
address to access, or what control action to perform. Here's a breakdown of what you're storing in your EX/MEM block and why each signal matters:

Stored Signals

Every signal and her purpose:

Signal	Purpose
PC+1	Used for instructions like JALR, where the return address (PC+1) may need to be written back to a register.
IMM	Forwarded in case the memory stage or later needs the immediate value again (especially for SW where the immediate is added to base to form address).
ALU forward (ALUFwr)	This is typically one of the source register values that may be used again in the MEM or WB stage, such as for a store instruction (e.g., SW R1, 0(R2) needs R1 forwarded).
ALU result (ALUres)	This is the result of the computation—either an arithmetic/logic result or a memory address (for LW/SW). It must be passed to MEM to read/write data or to later write back to the register file.
Destination Register (Rd)	Specifies which register will be written to in the Write Back stage (for instructions like ADD, ADDI, LW, etc.).
Operand B (B)	Often holds the value to store in memory for store instructions (e.g., SW).

Control Signals	These guide the MEM and WB stages on what to do—e.g., whether to read/write memory, write back to a register, or perform branching. Without passing these signals along, the next stages wouldn't know how to behave.
Jump	This signal may indicate whether the instruction is a jump, helping to manage control flow decisions and potential pipeline flushing.



In our pipelined processor, the necessary data and control signals are passed from the EX-stage to the MEM stage via the EX/MEM pipeline register. After

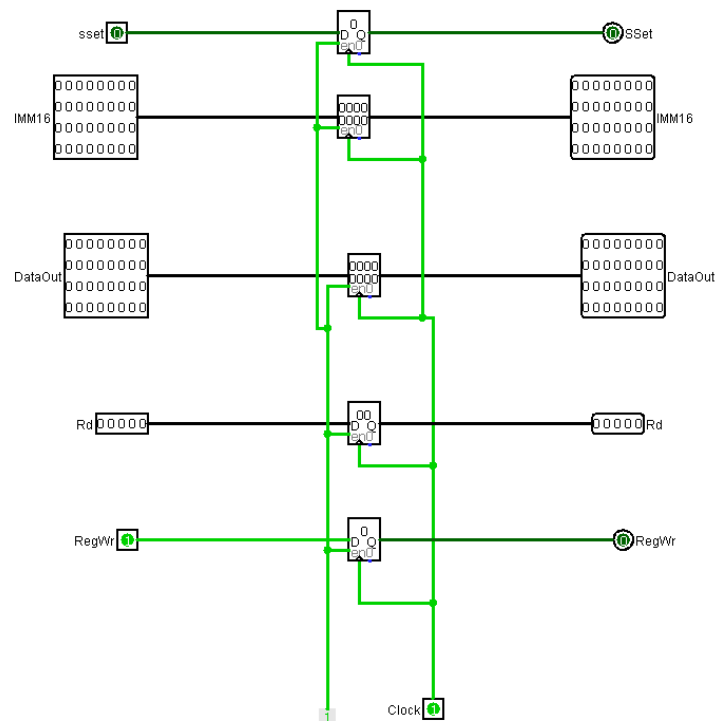
performing any required memory operations, the results—along with relevant metadata—are stored in the MEM/WB pipeline register, so the next stage (Write Back) can properly complete the instruction.

Operations Performed in the MEM Stage

- For load instructions (LW), the memory address is computed in the EX-stage and used here to read from data memory. The value retrieved is stored in the Data Out signal for later writing to the destination register.
- For store instructions (SW), the value to be written (from source register RS2) is provided via the EX-stage and written to the computed address in memory. These instructions typically don't require Write Back, so RegWrite is set to 0.
- For all non-memory instructions, this stage doesn't alter memory but still forwards important signals for the WB stage, like ALU result, IMM, or control flags such as SSET.

Stored Signals

Signal	Purpose in WB Stage
IMM16	Holds the immediate value used in some instructions (SET, SSET, etc.). This may be used in the Write Back stage if the instruction involves loading a constant.
Data Out	For LW, this is the value read from memory that will be written back to the destination register in WB.
Rd	The destination register number, indicating where the result (from ALU or memory) should be written.
RegWrite	A control signal that tells the Write Back stage whether it should write a value to the register file or not.
SSET	A control signal used for the SSET instruction. It tells the Write Back stage that it needs to perform the special concatenation operation instead of a normal write.



Write Back

The Write Back stage is the final stage in our pipelined processor. Its main purpose is to store the final result of an instruction into the register file, completing the instruction's execution. After the ALU computations, memory reads, or immediate value generation are done in previous stages, the result is written back to the destination register specified by the instruction.

This stage plays a crucial role in ensuring that the output of every instruction reaches the register file correctly and at the right time, especially since the pipeline allows multiple instructions to be in progress simultaneously.

Operations in the WB Stage

Depending on the instruction type, the value to be written back comes from different sources:

- For arithmetic and logic instructions, the result comes from the ALU output.
- For load instructions (LW), the value comes from Data Memory (stored in Data Out).
- For immediate instructions like SET or SSET, the value may come from the extended immediate field (IMM16), and may require special handling (as in the case of SSET).

The decision of whether to write back or not, and what value to write, is controlled by signals passed from the MEM stage—especially RegWrite and SSET.

Hazard Detect Forward and Stall unit

In a pipelined processor, multiple instructions are in different stages of execution at the same time. While this improves performance, it also introduces hazards—situations where the normal flow of instructions can lead to incorrect behavior. To maintain correct operation, our design includes a Hazard Detection Unit, a Forwarding Unit, and a Stall Mechanism that together resolve these hazards efficiently.



1. Data Hazards and Forwarding

A data hazard occurs when an instruction depends on the result of a previous instruction that hasn't completed yet. For example:

ADD R1, R2, R3

SUB R4, R1, R5 => R1 is not ready yet

Here, the SUB instruction needs the result of R1, but R1 hasn't been written back yet.

To solve this without stalling the entire pipeline, we use a Forwarding Unit (also called a Bypass Unit). This unit checks for dependencies by comparing the source registers of the current instruction in the EX-stage with the destination registers of instructions in the MEM or WB stage. If a match is

found and the result is ready, the data is forwarded directly from a later pipeline stage back to the EX-stage before the register file is updated.

Benefits:

- Avoids unnecessary stalls.
 - Improves pipeline performance by keeping instructions moving.
-

2. Load-Use Hazard and Stalling

Forwarding can't solve every hazard. A special case is the load-use hazard, where an instruction immediately uses data being loaded from memory. For example:

LW R1, 0(R2)

ADD R3, R1, R4 => R1 is not available yet

In this case, the data being loaded isn't ready until the MEM stage, so even with forwarding, the value isn't available in time for the ADD instruction in the next cycle.

To handle this, our Hazard Detection Unit detects this specific situation and inserts a stall:

- The pipeline is paused for one cycle.
- The ADD instruction is held in the ID stage, and a "bubble" (NOP) is inserted into the EX-stage.

- This gives the LW instruction time to complete its memory access and forward the result.

Benefits:

- Checks if the destination register of a LW in EX matches a source register of the instruction in ID.
 - Issues a stall signal to freeze PC and IF/ID for one cycle.
 - Allows safe continuation of dependent instructions.
-

3. Control Hazards

Though not part of the forwarding unit, it's worth noting that control hazards (from branches and jumps) are handled by:

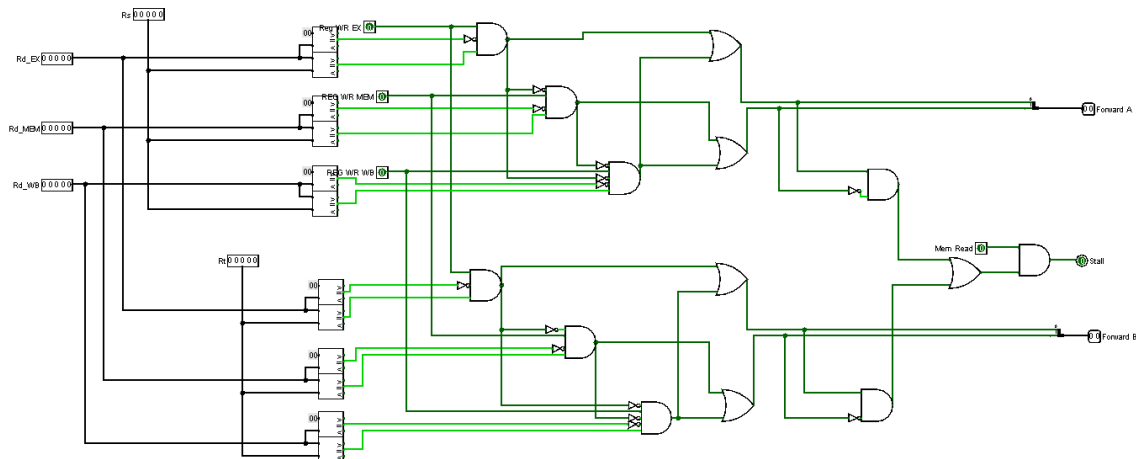
- Stalling the pipeline until the branch decision is known.
- Or optionally flushing incorrectly fetched instructions if the branch is taken.
- A branch predictor (bonus) can be added to minimize control stalls.

To improve performance in a pipelined processor, data hazards must be managed efficiently. The Forwarding Unit and Hazard Detection Unit are used together to avoid stalls when possible and insert them only when necessary.

Forwarding Unit

The Forwarding Unit handles data hazards by forwarding values directly from later pipeline stages (EX/MEM or MEM/WB) back to the ALU inputs, avoiding the need to wait for the write-back to the register file

Condition	Forward A	Forward B	Action
No Hazard	00	00	Use register file values
EX Hazard (Rs = Rd2 & EX. RegWrite)	10	00	Forward from EX/MEM to ALU (Rs)
EX Hazard (Rt = Rd2 & EX. RegWrite)	00	10	Forward from EX/MEM to ALU (Rt)
MEM Hazard (Rs = Rd3 & MEM.RegWrite)	01	00	Forward from MEM/WB to ALU (Rs)
MEM Hazard (Rt = Rd3 & MEM.RegWrite)	00	01	Forward from MEM/WB to ALU (Rt)
WB Hazard (Rs = Rd4 & WB. RegWrite)	11	00	Forward from WB to ALU (Rs)
WB Hazard (Rt = Rd4 & WB. RegWrite)	00	11	Forward from WB to ALU (Rt)
Multiple Hazards (Rs = Rd2 and Rd3)	10	00	Priority to EX/MEM (newer result)
Multiple Hazards (Rt = Rd3 and Rd4)	00	01	Priority to MEM/WB over WB



Forward A / Forward B: 2-bit control signals to select ALU source:

- 00: from register file
- 01: from MEM/WB
- 10: from EX/MEM
- 11: from WB stage

Hazard Detection Unit

When forwarding is **not enough** (especially in **load-use hazards**), we need to **stall** the pipeline to prevent incorrect execution.

Load-Use Hazard:

Occurs when an instruction in EX depends on the result of a **load instruction** that is still in ID/EX.

Condition:

if (ID/EX.MemRead &&

((ID/EX. RegisterRt == IF/ID. RegisterRs) || (ID/EX. RegisterRt == IF/ID.

RegisterRt)))

Actions:

- Stall IF/ID and ID/EX (freeze PC and pipeline registers)
- Insert **NOP** in EX stage
- Disable writes to PC and IF/ID

Condition	ForwardA	ForwardB	MemToReg	Stall
forward from EX/MEM (ForwardA = 10, Forward B = 10), but it's a load (MemToReg = 1)	10	10	1	1

Load Delay

In pipelined CPUs, forwarding (bypassing) resolves data hazards by sending results directly from later stages to earlier ones. A "delayed forward" scenario might involve:

Problem: A result isn't ready in time for forwarding (e.g., multi-cycle operations).

Solution: Insert a pipeline stall until the data is available, then forward it.

PC Control Unit

It controls when to allow or prevent updates to the Program Counter (PC) — essential for handling branches, jumps, and pipeline hazards like flushes (kills).



Inputs:

- Branch: High when a branch instruction is active.
- Jump: High when a jump instruction is active.

Outputs:

- PC Control Unit: Controls whether the PC is updated. (0 = allow update, 1 = stall)
- Kill 1: Tells the pipeline to **flush/kill** the instruction in the IF stage.


Branch	Jump	PC Control Unit	Kill 1
0	0	1 (stall)	0
0	1	1 (stall)	1 (flush)
1	0	1 (stall)	1 (flush)
1	1	0 (allow update)	1 (flush)

- **Kill 1 = Branch OR Jump**

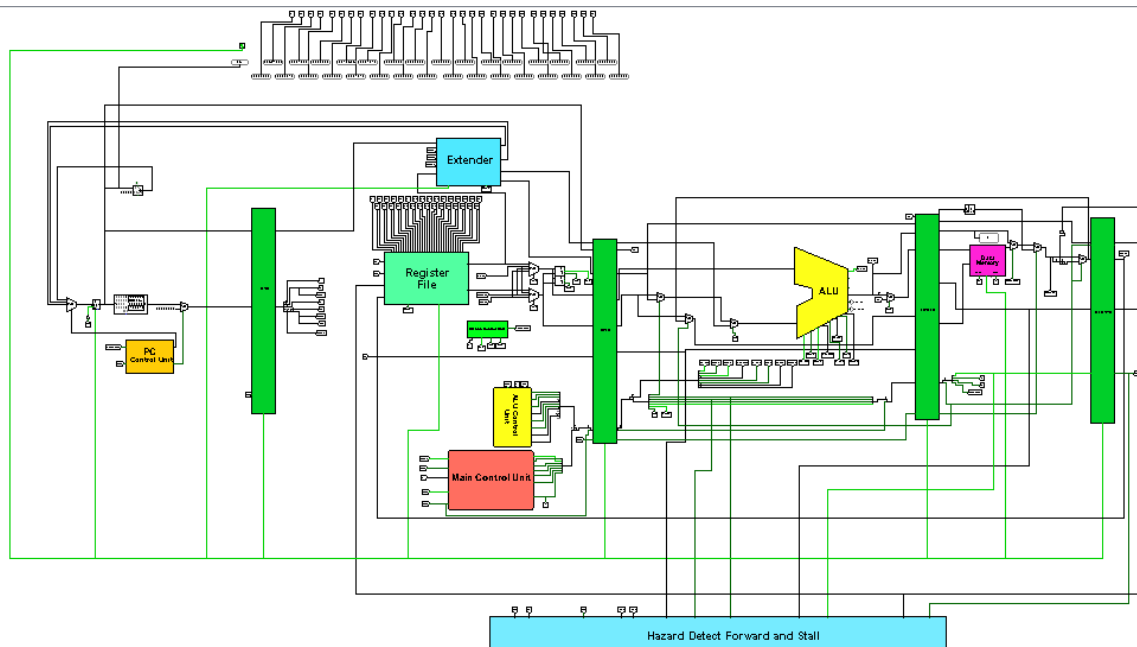
→ So it's 1 when **either** Branch or Jump is 1 (flush the fetched instruction).

- **PC Control Unit = NOT (Branch AND Jump)**

→ Only allows PC update (0) when **both** Branch and Jump are 1. Otherwise, it **stalls** the PC (1 = freeze).



Full view of Pipelined Processor



Finally, our pipelined processor is complete and fully functional! We've built all five stages — Fetch, Decode, Execute, Memory, and Write Back — and connected them with proper pipeline registers.

We added forwarding and hazard detection units to make sure the processor handles data hazards without unnecessary stalling. For control hazards like branches and jumps, we used a Kill signal and PC Control logic to flush or freeze the pipeline when needed.

After testing everything together, we're proud to say that our processor runs smoothly and handles all kinds of instructions correctly. It was a challenging project, but seeing the whole pipeline working in sync was definitely worth the effort!

Team Work

Omar Ahmed Othman

- Worked more than 90 hours on Design and Implementation.
- Designed ALU Control Unit, Register File, PC, Main Control Unit, Pipelined Registers Stage, Extender.
- Put All components together and formed the full view of the Single Cycle Processor and the full view of Pipelined Processor
- Did Simulation and Testing and solved some bugs during execution.
- Updated the Project Report and enhanced details for both phases.

Ahmed Osama Abd-Elghaffar

- Worked more than 80 hours on Design and Implementation.
- Designed ALU 1-bit, ALU 32-bit, Branch Control Unit, Main Control Unit, PC Control Unit for Pipeline, Hazard Detect Forward and Stall Unit.
- Helped in putting All components together in Single Cycle and Pipelined Processor
- Did Simulation and Testing and solved some bugs during execution.
- Updated Register Stages and Solved Bugs during Execution in Pipeline.

Mohamed Nady Mahmoud

- Worked more than 30 hours on Report.
- Designed Instruction Memory and Data Memory and helped in others stages in Pipeline
- Helped in putting all components together.

